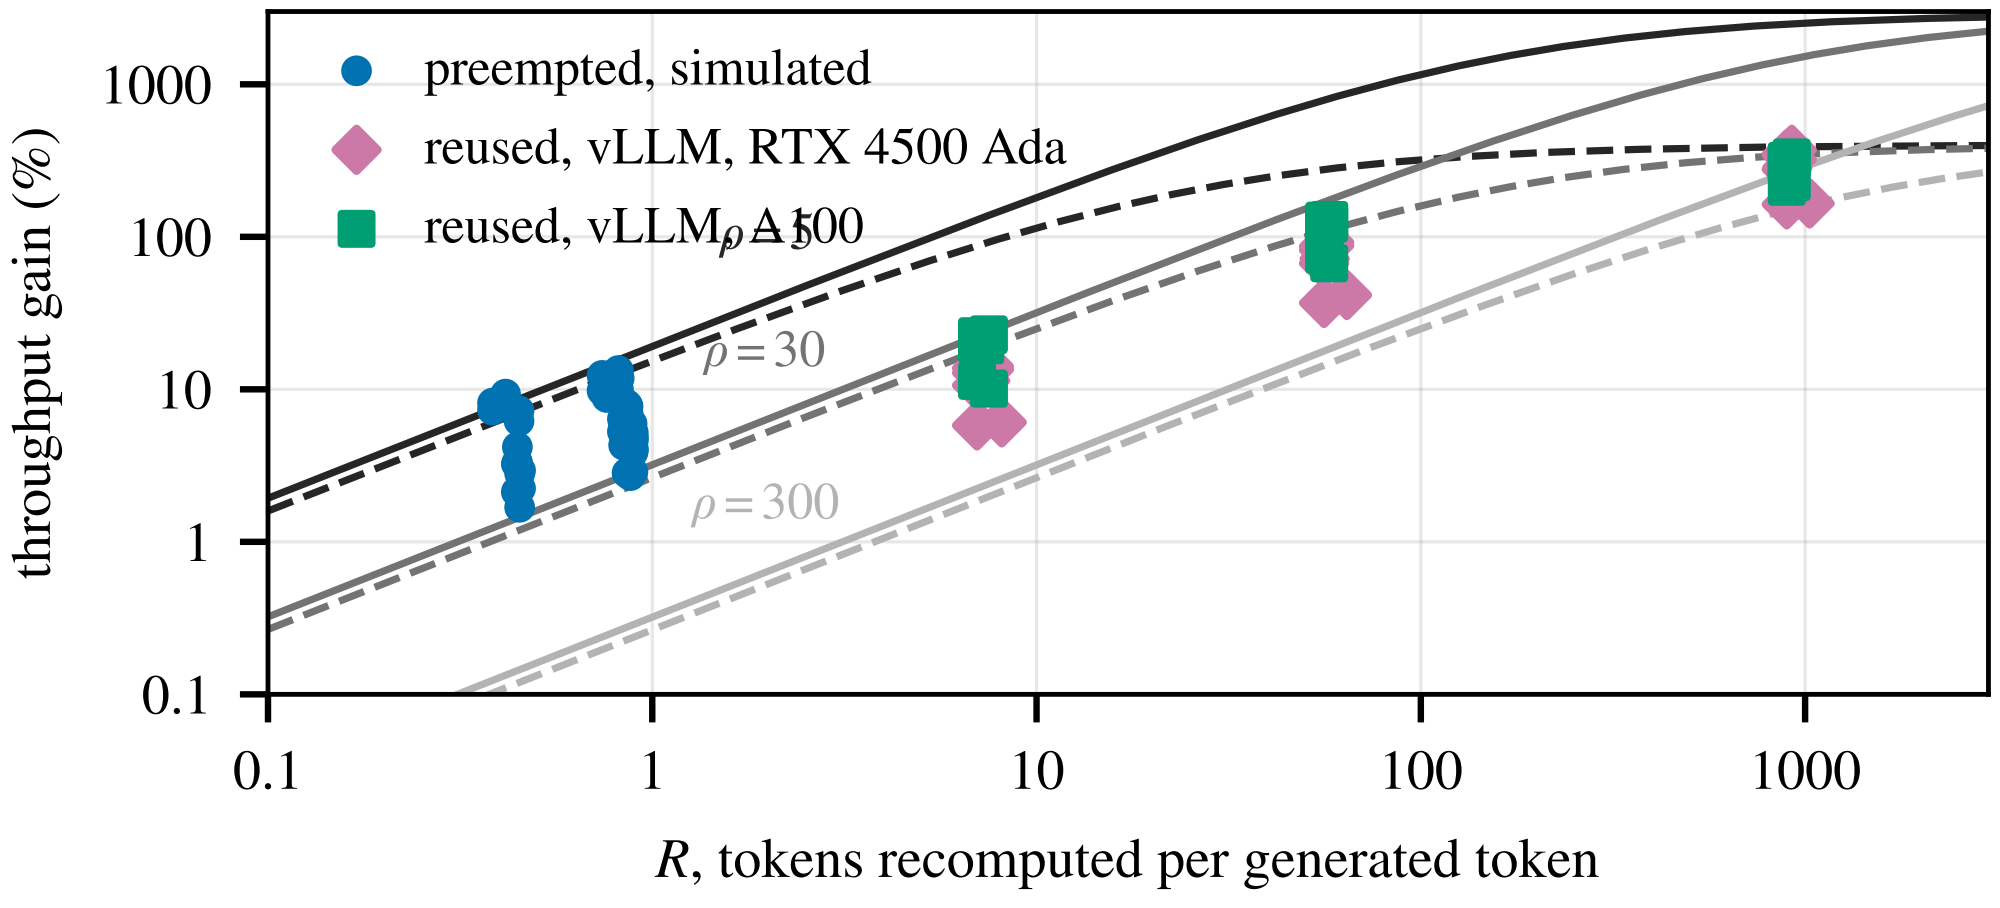
preempted, simulated
reused, vLLM, RTX 4500 Ada
reused, vLLM, A100
$\rho = 30$
$\rho = 300$
throughput gain (%)
1000
100
10
1
0.1
0.1
1
10
100
1000
R, tokens recomputed per generated token