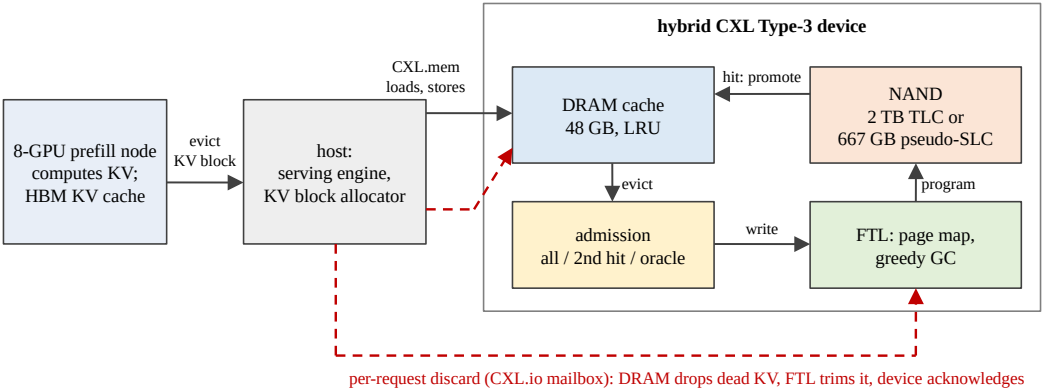
hybrid CXL Type-3 device
8-GPU prefill node
computes KV;
HBM KV cache
evict
KV block
host:
serving engine,
KV block allocator
CXL.mem
loads, stores
DRAM cache
48 GB, LRU
hit: promote
NAND
2 TB TLC or
667 GB pseudo-SLC
evict
program
admission
all / 2nd hit / oracle
write
FTL: page map,
greedy GC
per-request discard (CXL.io mailbox): DRAM drops dead KV, FTL trims it, device acknowledges